



UiT The Arctic University of Norway

Project 2: Non-preemptive Scheduling



INF-2201 Operating System Fundamentals

Spring 2023

Overview

- In P1 you implemented a boot loader
 - But the OS kernel was very minimalistic
- Here you implement a simple OS kernel
- You can (but are not required to) use the code you wrote in P1

Kernel

- Multitasking support
- *Non-preemptive* scheduling
- A running process/thread has to relinquish control explicitly by:
 - Yielding
 - Exiting
 - Blocking on a lock


Processes and Threads

- Processes, threads and the kernel share the same flat address space
 - But how do they differ?
 - And how do they differ from Linux Threads and Processes?
- No real protection
- Separate address spaces later

Protected Mode

- Kernel runs in protected mode
 - Setup done in boot loader
- Everything runs at highest CPU privilege level (ring 0)
- More about protected mode in colloquium groups

Tasks

- Initialization
 - Process Control Block (PCB)
 - Context switching
 - System call mechanism
 - Stacks
 - Synchronization
 - Assembly / inline assembly
 - Performance measurements
 - Design review
- 

Initialization

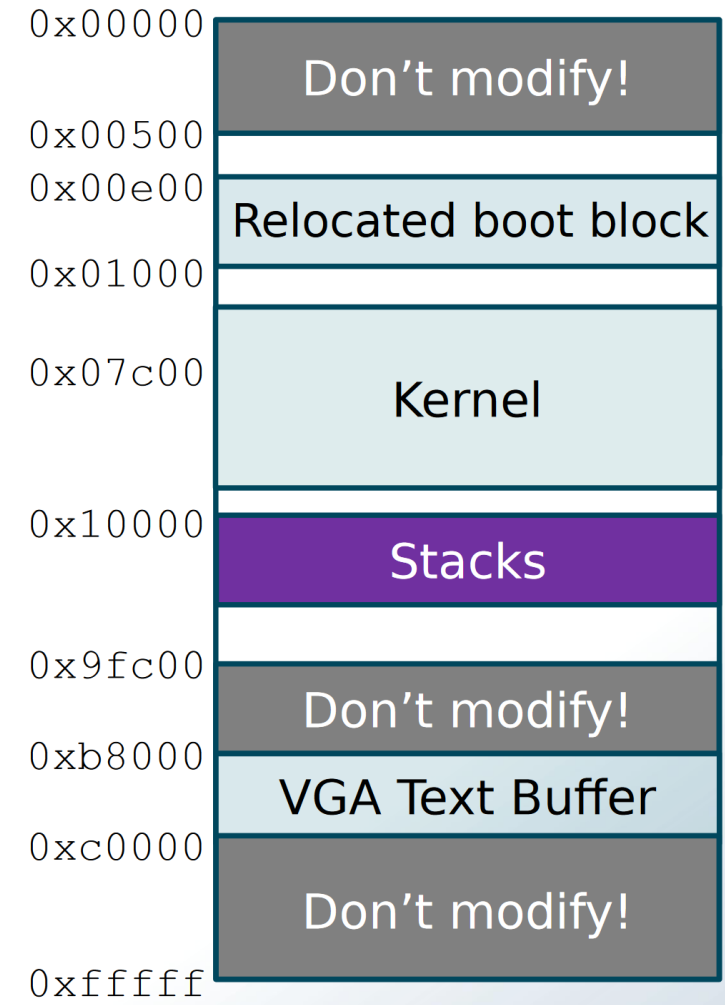
- Initialization
- Setup the required data structures for processes and threads
 - No dynamic loading
 - Everything set at startup

Process Control Block (PCB)

- In kernel.h
- What should be in the PCB?
 - pid?
 - is_thread?
 - stack?
 - next, previous PCB?
- Anything else?
 - What is in the Linux PCB? Or Windows PCB?
 - \Rightarrow Describe this in the design review!

Stacks

- How many stacks?
 - 2 per process, 1 per thread
 - Why?
- Where to put them in memory?
 - Upper limit: 0x9fc00
 - Suggestion: between 0x10000 and 0x20000
- Stack size?
 - 8KB should be fine.



Memory layout
(not to scale)

Scheduler

- Simple non-preemptive scheduling
 - IS round-robin good enough?
 - How to do blocking and unblocking?
- ⇒ Design review!

Context Switch Procedure

- How to switch between processes and threads?
 - They must call `yield()` explicitly (non-preemptive)
 - In next project: time slice expires (preemptive)
- What to save?
 - GPRs?
 - More?
- Where to save it?
 - Stack?
 - PCB?
- ⇒ Design review!


System Call Mechanism

- How does a process get services from the kernel?
 - This assignment: a special function call using a single entry “jump table”
 - Later: interrupt/trap mechanism
- Why cannot system calls be implemented as ordinary function calls?
- Is our approach better?

System Call Mechanism

- At runtime, load the address of `kernel_entry()` into memory location `0xf00`
 - Already done in pre-code
- Prototype: `void kernel_entry(int fn);`
- Define this in `syslib.h`
`#define ENTRY_POINT ((void (**)) 0xf00)`
 - Pointer to pointer to function with non-defined argument list returning void
 - ...at address `0xf00`
- Declare the following in `syslib.c`
`static void (**entry_point) () = ENTRY_POINT;`
- `entry_point` has address `0xf00`, and `*entry_point` is the address of our kernel entry point function
 - `*entry_point = kernel_entry;` (done in `kernel.c`)
- Macro for invoking syscall:
`#define SYSCALL(i) ((*entry_point)(i))`

kernel_entry() in memory




The kernel sets the memory at address `0xf00` to the address of the `kernel_entry()` function. The user process can then read the memory at `0xf00` to learn where the `kernel_entry()` function is located in memory.

Synchronization

- Locks are only used by threads
- Many threads can try to acquire a lock
 - Need to maintain queue of threads waiting for a lock
 - Where?
- `lock_acquire()`
 - Check lock
 - Got lock? Great!
 - If not, block itself
- `lock_init()`, `lock_release()`

Measure Context Switch Time


- What to measure and how (methodology)
 - \Rightarrow Design review
 - Do measurements
 - Get results
 - Report and discuss results in report
- 

Extra Credit 1 – time

- Implement something similar to the Unix command “time”
 - Measure time spent in user mode and kernel mode
- Get an A+ grade
- ..or a TA job next year?

Extra Credit 2 and 3 – more threads and processes

- Add a new thread to your kernel
- Add a new process to your OS



Design review

- For each assignment you are required to give a design report for you plan to implement the solution
 - Note! describe it at the “design level”
 - Not a formal presentation - just you and your TA
- You should be prepared
 - Have something to show during the design review
 - Only oral presentation is not acceptable
 - You need to convince the TA that you understand the project
- The design review is a mandatory “assignment”:
 - Pass / no-pass grading

Possible topics for design review

- What's in the PCB? Why is it there?
- How will you implement locks? What about process queues?
- Where should context be saved?
- What is the context (what must be saved?)
- What is the difference between processes and threads?
 - In this assignment?
 - How do “our” processes and threads differ from “normal” processes and threads?
- How many stacks per process? And why?
- How many stacks per thread? And why?
- Non-preemptive vs preemptive scheduling; What is similar? What is different?

Code

- Comment your code
 - Give a high-level overview of what the code does
 - Especially important when dealing with hardware, entry points, and synchronization/ blocking
 - Comments are part of the grading
- Test your code
- Don't cheat

Report

- Maximum 4 pages
- Give an overview of how you solved each task
 - including extra credits
- Describe how you have tested your code
 - and known bugs/ issues
- Describe the methodology, results, and conclusions for your performance measurements

Handin

- This is an exam, so there is a strict deadline
 - Extensions only for valid and preapproved reasons
- Wiseflow for handin
 - PDF of your report
 - Zip file with source code (remember to make clean before zipping)
- Add your name, GitHub username, and email to the report

Hints - Flat Address Space

- The bootblock code:
 - Switches to protected mode
 - Sets up the CS, DS, SS, and extra segment registers
 - You can use the entire memory by utilizing registers like EAX
 - Do NOT modify the segment registers
 - You have access to the first 1MB of memory
 - Actually, you have access to a 32-bit address space, but we won't use memory > 1MB limit in this assignment)
 - Including the video-memory area (VGA text buffer @ 0xb8000)
 - Make sure your code, data, and stacks do not overwrite "important" memory areas (see the memory map)

Hints - Threads and Processes

- Your OS will have processes and threads
 - Threads are linked with the kernel
 - Processes are separate executable files
 - Note: they all end up together in the same image file
- Processes need a special way to call functions in the kernel
- Threads can just call kernel functions directly
- In a more advanced OS, there are many other differences between threads in the kernel and processes

Hints – Synchronization

- The synchronization primitives can only be used by threads within the kernel
- The two functions `block()` and `unblock()` are to be used by the synchronization code
- Blocking a process is not specific to locks, but is a general purpose service that the kernel should support