

UiT

THE ARCTIC
UNIVERSITY
OF NORWAY

Project 3

Preemptive Scheduling

INF-2201 Operating System Fundamentals
Spring 2023

Department of Computer Science
University of Tromsø



Overview

- You will implement an OS that schedules threads and processes *preemptively*.
- Implement synchronization primitives that work with preemption
 - Re-implement locks
 - Condition variables
 - Semaphores
 - Reusable barriers
- Concurrent programming
 - Implement a fair solution to the dining philosophers problem
- Extra credits
 - Priority scheduling
 - Pthreads implementation of dining philosophers

Previous operating system

- **Non-preemptive** multiprogramming kernel
- System calls by **calling a function in the kernel**
- Processor executes in protected mode
 - All processes run in ring 0
- Synchronization (**locks**) only for kernel threads
 - For non-preemptive kernel

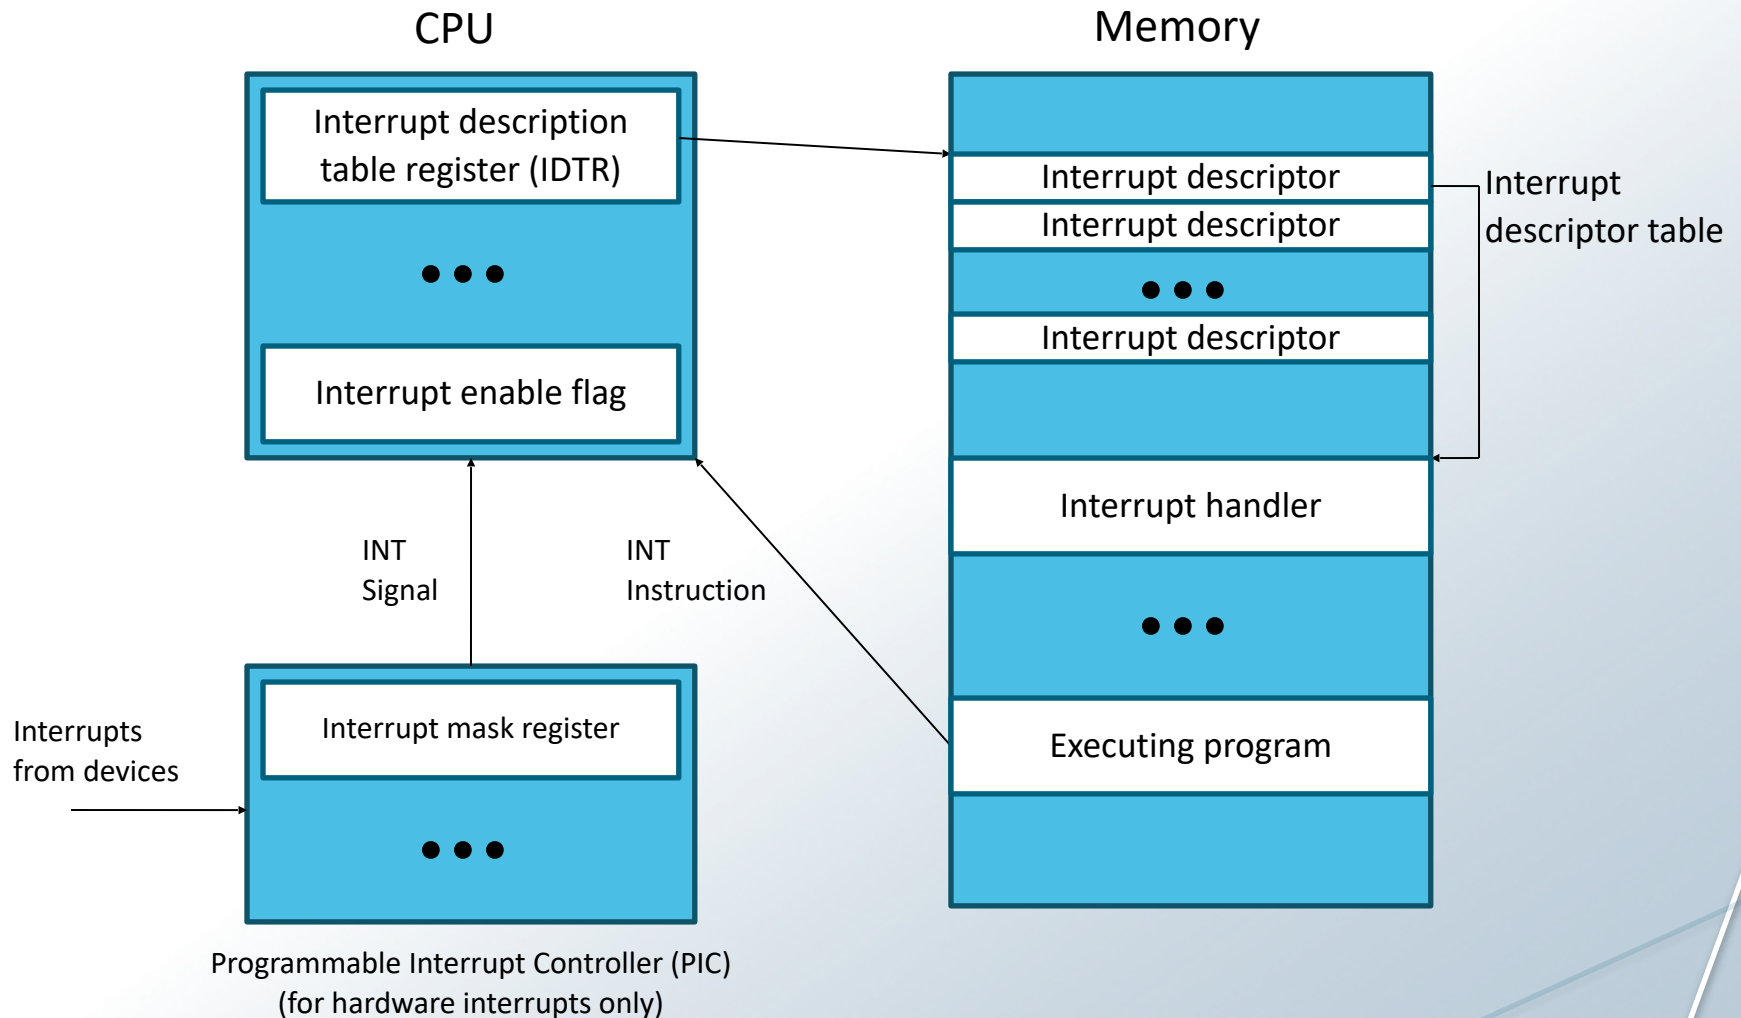
New operating system

- **Preemptive** multiprogramming kernel
- System calls via **software interrupts**
- Processor executes in protected mode
 - Processes will still run in ring 0
- Synchronization primitives that **work with preemption** (but still only for kernel threads)
 - Locks
 - Semaphores
 - Condition variables
 - Barriers

Three types of interrupts

- Hardware interrupts (external interrupts)
 - Example: system timer interrupt
- Software interrupts (INT instructions)
 - Example: (real-mode) BIOS calls
- Software exceptions

Hardware support for interrupts



What does the CPU do when an interrupt occurs from ring 0?

- Get interrupt descriptor address
- Make privilege checks
- Push `EFLAGS`, `CS`, and `EIP` on stack
- Clear `EFLAGS[IF]` (and `EFLAGS[TF]`)
- Load `CS:EIP` from interrupt descriptor
 - Jumps to interrupt handler

And then?

- The interrupt handler executes.
- If interrupt came from PIC (if it was a hardware interrupt), the handler must send an end-of-interrupt (`EOI`) signal to the PIC.
 - Because PIC sets a bit in its In-Service Register (`ISR`), corresponding to the hardware interrupt vector, to block consecutive interrupts on that vector until handler completes.
 - `EOI` clears interrupt vector bit in `ISR`.
 - In other words, if you forget to clear `EOI`, you will only get one interrupt from each hardware vector ...ever!
- To return from interrupt handler, execute `IRET` instruction
 - Pops `CS`, `EIP` and `EFLAGS`.
 - When restoring `CS:EIP`

Pseudo-C template of an interrupt handler

```
void irqX(void)
{
    /* save context of interrupted thread/process */
    ...
    /* do the work */
    ...
    /* restore context of interrupted thread/process */
    ...
    /* restore EIP, CS and EFLAGS */
    asm volatile("iret");
}
```

Remember: if this is a H/W interrupt handler, you also need to issue EOI:

```
outb(0x20, 0x20);
```

Preemptive Scheduling

- Most of it is already setup for you
- Setup system timer interrupt so that an interrupt will be generated on IRQ line 0 every 10 milliseconds (given in `kernel.c`).
 - You may want to change this value for debugging and testing.
- You need to implement the code for the timer interrupt handler in `entry.S`.
 - Switch between user and kernel stack
 - Save and restore the context
- Modify `yield()`, `lock_acquire()`, `lock_release()`, etc. to deal with preemptive scheduling.
 - Critical regions (`enter_critical()`, `leave_critical()`) in `entry.S`

Synchronization primitives

- Semaphores, barriers, and condition variables with semantics as discussed in lectures
 - Also check how this is done in Pthreads
 - The template for code you need to write is in `thread.c`.
- You will also need to re-implement locks (like in P2), but this time make sure that they can be used in a preemptive context (avoid race conditions).
- You have to design data structures for semaphores, barriers and condition variables.
 - Empty structures are defined in `thread.h`.

Dining philosophers



- Precode comes with a solution using semaphores
 - Three philosophers: Caps, Num and Scroll
 - Watch the keyboard LEDs.
 - But this solution is not fair.
 - Favors Caps (why?)
- Come up with and implement a solution that is “fair”
 - Every philosopher can eat for the same amount of time
- Document your solution

Barriers (1)

- A barrier is a synchronization mechanism where several threads can be “realigned”, that is, once they leave the barrier, they will all be in the same stage of execution.
- An n-barrier works by blocking threads calling `barrier_wait()`, then unblocks all the threads when the *n*th thread calls `barrier_wait()`.

Barriers (2)

0: `barrier_init(3)`



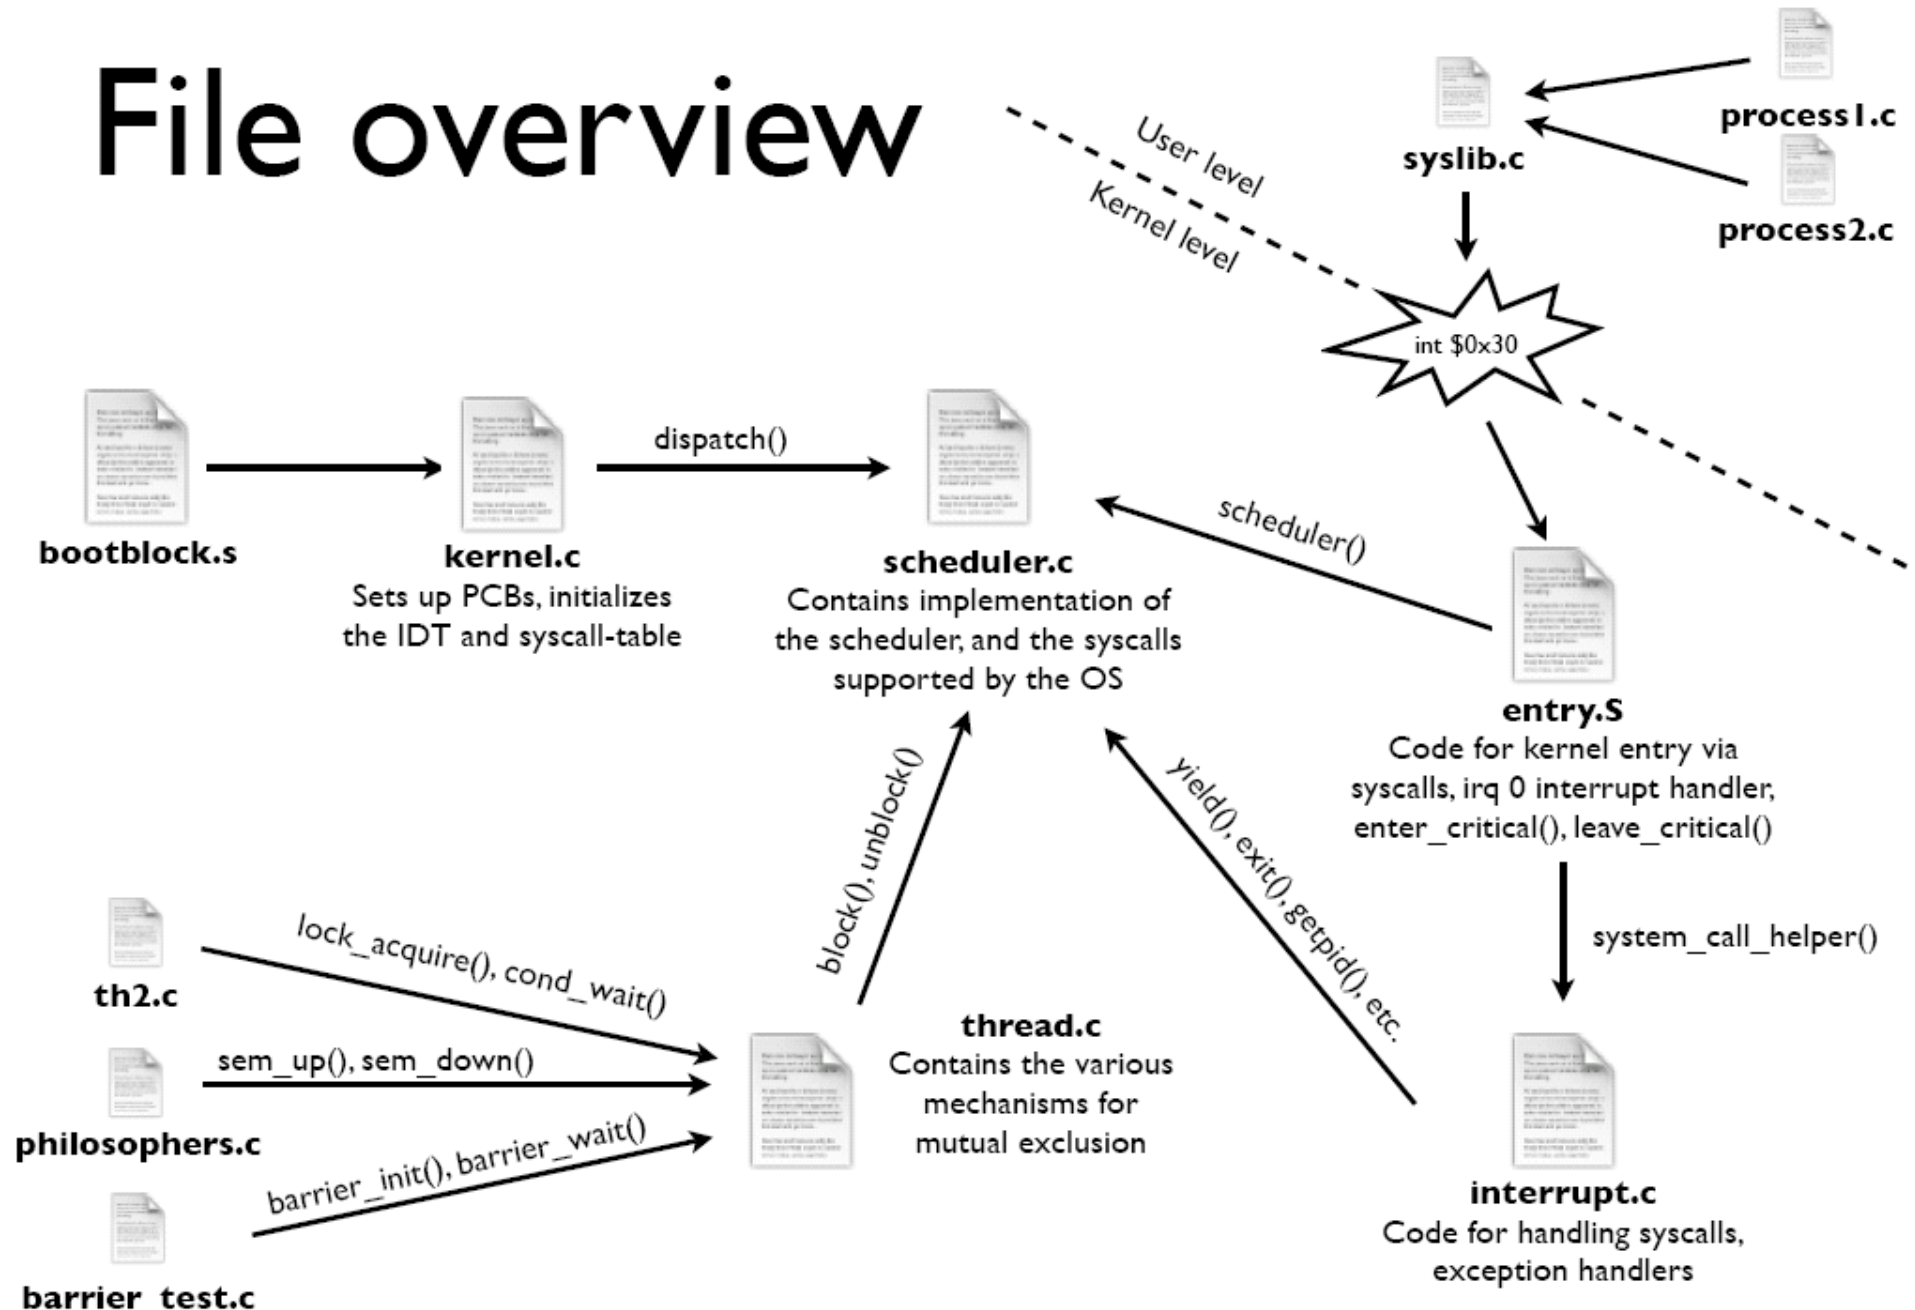
Barriers (3)

- Beware of race conditions!
 - What happens when one thread manages to call `barrier_wait()` a second time before all the other threads have been unblocked

Extra Credits

- Priority scheduling
 - Process 1 calls `set_priority()`
 - Modify `scheduler()`
- Dining philosophers implemented using Pthreads
 - You will do this in Linux!
 - You need to create the source files and Makefile (there is no pre-code)
 - Refer to online documentation for Pthreads details

File overview



Where to start?

- Do the preemptive-multitasking bit first!
- Once that is done...
 - Complete the locks
 - Semaphores
 - Condition variables
 - Barriers
- Remember, all synchronization primitives must work with preemption present!
- Finally, work on the dining philosophers problem
- You really should learn how to use these mechanism in Pthreads

Design review, code, report, hand-in

- Same as for P1
- ...using GitHub